

# **PPDS Week 5 Wednesday**

**Pandas continued**

**Data science with python**

Karl-Emil Kjær Bilstrup

17.12.25

# Python REPL Advertisement

## Read, Evaluate, Print, and Loop

For exploration and to test assumptions

```
% python
>>> pd_series = pd.Series({'a':1.0, 'b':2.0, 'c':3.0})
>>> pd_series
a    1.0
b    2.0
c    3.0
dtype: float64
```

# Creating a Series

```
import pandas as pd

# From a list with custom index
pd_series = pd.Series([1.0, 2.0, 3.0],
                      index=['a', 'b', 'c'])

# From a dictionary
pd_series = pd.Series({'a':1.0, 'b':2.0, 'c':3.0})

print(pd_series)
# Output:
# a    1.0
# b    2.0
# c    3.0
```

# Creating a DataFrame

```
import pandas as pd
import numpy as np

# From a dictionary of lists
data = {
    'Name': ['Alice', 'Bob', 'Charlie'],
    'Age': [25, 30, 35],
    'City': ['NY', 'SF', 'LA']
}
df = pd.DataFrame(data)

print(df)
#      Name  Age City
# 0   Alice   25   NY
# 1     Bob   30    SF
# 2 Charlie   35    LA
```

# Basic DataFrame Operations

```
# View first few rows
print(df.head())

# Get column names
print(df.columns)

# Access a column
print(df['Name'])

# Add a new column
df['Country'] = ['USA', 'USA', 'USA']

# Basic statistics for numeric columns
print(df.describe())
```

# Reading and Writing Data (I/O)

Pandas makes it easy to work with different file formats:

	Read	Write
general	read_table()	
csv	read_csv()	to_csv()
html	read_html()	to_html()
json	read_json()	to_json()
excel	read_excel()	to_excel()
sql	read_sql()	to_sql()
...		

# pandas - Dataframe - methods

Numpy functions apply to dataframes

```
np_array = np.arange(9).reshape((3,3))
df = pd.DataFrame(np_array, index=['a', 'b', 'c'],
                  columns=['col1', 'col2','col3'])

# stats
print(df.mean(axis=0), df.std(axis=0), df.median(axis=0))

# matrix multiplication
print(np.dot(df, df.T))
```

axis: 0 or index -> along index; 1 or columns -> along columns

# Basic Pandas DataFrame Methods

Method	Description	Example
<code>head()</code> , <code>tail()</code>	View first/last rows	<code>df.head(3)</code>
<code>info()</code>	Get DataFrame info (types, memory)	<code>df.info()</code>
<code>describe()</code>	Summary statistics	<code>df.describe()</code>
<code>shape</code>	Get dimensions (rows, columns)	<code>df.shape</code>
<code>columns</code>	Get column names	<code>df.columns</code>
<code>dtypes</code>	Get data types of columns	<code>df.dtypes</code>
<code>value_counts()</code>	Count unique values	<code>df['column'].value_counts()</code>
<code>sort_values()</code>	Sort by values	<code>df.sort_values('column')</code>
<code>groupby()</code>	Group data	<code>df.groupby('column').mean()</code>
<code>dropna()</code>	Remove missing values	<code>df.dropna()</code>
<code>fillna()</code>	Fill missing values	<code>df.fillna(0)</code>
<code>reset_index()</code>	Reset index to default	<code>df.reset_index()</code>

# Common Selection Methods

Operation	Description	Example
Select column	Get single column	<code>df['column']</code>
Select multiple columns	Get several columns	<code>df[['col1', 'col2']]</code>
Select by label	Get rows by label	<code>df.loc[row_label]</code>
Select by position	Get rows by position	<code>df.iloc[0:2]</code>
Boolean indexing	Filter with condition	<code>df[df['Age'] &gt; 30]</code>
Set column	Add/modify column	<code>df['new_col'] = values</code>

Pandas provides specific support for Series of certain types, through so-called Accessors.

Data type	Accessor
String	str
Categorical	cat
Datetime	dt
Sparse	sparse

```
pd_series.str.upper()
```

# pandas - Dataframe - selecting with boolean array

```
np_array = np.arange(9).reshape((3,3))
df = pd.DataFrame(np_array, index=['a', 'b', 'c'],
                  columns=['col1', 'col2','col3'])
# I want to take all values greater than 4
mask = (df > 4)
print(df[mask])
    col1  col2  col3
a    NaN    NaN    NaN
b    NaN    NaN    5.0
c  6.0  7.0  8.0
# I want to zero out all large values
df[mask] = 0
print(df)
    col1  col2  col3
a    0    1    2
b    3    4    0
c    0    0    0
```

# pandas - Dataframe - selecting with boolean array

You can also create masks for selecting whole rows or columns

Selecting rows:

```
np_array = np.arange(9).reshape((3,3))
df = pd.DataFrame(np_array, index=['a', 'b', 'c'],
                  columns=['col1', 'col2','col3'])

# take all rows whose first element is greater than 2
row_mask = df.iloc[:,0] > 2
print(df[row_mask])
# or print(df.loc[row_mask])
   col1  col2  col3
b    3     4     5
c    6     7     8
```

# pandas - Dataframe - selecting with boolean array

Selecting columns:

```
np_array = np.arange(9).reshape((3,3))
df = pd.DataFrame(np_array, index=['a', 'b', 'c'],
                  columns=['col1', 'col2','col3'])

# take all columns whose elements sum to greater 10
col_mask = df.sum(axis=0) > 10
df.loc[:,col_mask]
col2  col3
a      1    2
b      4    5
c      7    8
```

# Exercise 1

1. Download the `british-english` file from absalon or Github
2. Read it into a `DataFrame`. Are there any options you need to add to make it work?  
**Hint 1: Print the dataframe and see if you can find an issue. Hint 2: Check the Header parameter in the [documentation](#). Hint 3: Continue to 3. and see if it helps you spot the issue.**
3. Figure out how to assign a different column name - e.g. 'words'  
**Hint: Start by figuring out how to print the column name, then assign a new name.**
4. Figure out how to select all rows starting with the letter A.  
**Hint 1: Use the string accessor on a column. Hint 2: See `str.startswith` methods in [documentation](#).**

# pandas - Dataframe - adding columns

Add a new column by assigning to the relevant label

```
np_array = np.arange(9).reshape((3,3))
df = pd.DataFrame(np_array, index=['a', 'b', 'c'],
                  columns=['col1', 'col2','col3'])
# Adding a copy of col1 to col4
df.loc[:, 'col4'] = [0,3,6]
# Using df.assign
df = df.assign(col4 = df.loc[:, 'col1'])
# Using df.insert
df.insert(3, 'col4', df.loc[:, 'col1'])
print(df)
   col1  col2  col3  col4
a      0      1      2      0
b      3      4      5      3
c      6      7      8      6
```

# pandas - Dataframe - adding columns

If columns don't have the expected index, missing elements will be set to NaN

```
np_array = np.arange(9).reshape((3,3))
df = pd.DataFrame(np_array, index=['a', 'b', 'c'],
                  columns=['col1', 'col2','col3'])
# Adding a copy of col1 to col4
df.loc[:, 'col4'] = df.loc['b':'', 'col1']
print(df)
   col1  col2  col3  col4
a      0      1      2    NaN
b      3      4      5      3
c      6      7      8      6
```

# pandas - Dataframe - adding rows

Add a new row by assigning to the relevant label

```
np_array = np.arange(9).reshape((3,3))
df = pd.DataFrame(np_array, index=['a', 'b', 'c'],
                  columns=['col1', 'col2','col3'])
# Adding a copy of a to d
df.loc['d',:] = df.loc['a',:] # or [0,1,2]
# df.insert and df.assign only support adding columns
print(df)
   col1  col2  col3
a    0    1    2
b    3    4    5
c    6    7    8
d    0    1    2
```

# pandas - Dataframe - merging

`pd.merge()` combines two dataframes based on a shared column label

```
df1 = pd.DataFrame({'id' : [1, 2, 3],  
'gender' : ['male', 'female', 'female']})  
  
df2 = pd.DataFrame({'id' : [1, 2, 3],  
'name' : ['Bob', 'Alice', 'Anna']})  
  
df = pd.merge(df1, df2, on='id')  
print(df)  
  
      id  gender    name  
0     1    male     Bob  
1     2   female   Alice  
2     3   female   Anna
```

# pandas - Dataframe - merging

You can also merge on the index labels

```
df1 = pd.DataFrame({'id' : [1, 2, 3],  
'gender' : ['male', 'female', 'female']})  
df2 = pd.DataFrame({'id' : [1, 2, 3],  
'name' : ['Bob', 'Alice', 'Anna']})  
df = pd.merge(df1, df2, left_index=True, right_index=True)  
print(df)
```

	id_x	gender	id_y	name
0	1	male	1	Bob
1	2	female	2	Alice
2	3	female	3	Anna

# pandas - Dataframe - concatenation

`pd.concat()` combines multiple dataframes along columns or rows, controlled by the argument `axis`

```
np_array = np.arange(6).reshape((2,3))
df1 = pd.DataFrame(np_array, index=['a', 'b'],
                    columns=['col1', 'col2', 'col3'])
df2 = pd.DataFrame(np_array, index=['d', 'e'],
                    columns=['col1', 'col2', 'col3'])

# Concatenation along index
df = pd.concat([df1, df2], axis='index') # axis: {0/'index', 1/'columns'}, default 0
print(df)
   col1  col2  col3
a    0    1    2
b    3    4    5
d    0    1    2
e    3    4    5
```

# pandas - Dataframe - concatenation

`pd.concat()` combines multiple dataframes along columns or rows, controlled by the argument `axis`

```
np_array = np.arange(9).reshape((3,3))
df1 = pd.DataFrame(np_array, index=['a', 'b', 'c'],
                    columns=['col1', 'col2','col3'])
df2 = pd.DataFrame(np_array, index=['a', 'b', 'c'],
                    columns=['col4', 'col5','col6'])

# Concatenation along index
df = pd.concat([df1,df2],axis='columns') # axis: {0/'index', 1/'columns'}, default 0
print(df)
   col1  col2  col3  col4  col5  col6
a      0      1      2      0      1      2
b      3      4      5      3      4      5
c      6      7      8      6      7      8
```

# pandas - Dataframe - groupby

`df.groupby()` splits the dataframe `df` into different groups

the keyword `by=` determines the grouping rule

```
df = pd.DataFrame({'Animal': ['Falcon', 'Falcon',
                               'Parrot', 'Parrot'],
                   'Max Speed': [380., 370., 24., 26.]})
# I want to calculate the average speed of both species
groups = df.groupby(['Animal'])
# Apply mean to every animal group
avg_speed_df = groups.mean()
print(avg_speed_df)
```

	Max Speed
Animal	
Falcon	375.0
Parrot	25.0

# pandas - DataFrame - groupby

You can loop over the returned variable to check what the groups are

```
df = pd.DataFrame({'Animal': ['Falcon', 'Falcon',
                               'Parrot', 'Parrot'],
                   'Type': ['Captive', 'Wild', 'Captive', 'Wild'],
                   'Max Speed': [380., 370., 24., 26.]})
# I want to calculate the average speed of both species
groups = df.groupby(['Animal','Type'])
# I want to check what the groups are
for name, value in groups:
    print(name,value)
('Falcon', 'Captive')      Animal  Type  Max Speed
                           0   Falcon  Captive    380.0
('Falcon', 'Wild')        Animal  Type  Max Speed
                           1   Falcon   Wild     370.0
# ... prints two more groups
```

# pandas - DataFrame - sorting values

`df.sort_values(by=label)` sorts either by row or column values:

```
np_array = np.random.rand(3,3)
df = pd.DataFrame(np_array,
                  index=['a', 'b', 'c'],
                  columns=['col1', 'col2', 'col3'])

print(df)
print(df.sort_values(by='col2', ascending=False))
print(df.sort_values(by='a', axis='columns'))
```

# pandas - Remarks

- Indexing and slicing - return type:
  - `[ :, : ]` -> dataframe
  - `[ :, i ] , [ i , : ]` -> series
  - `[ i , j ]` -> single element
- pandas functions/methods are not in-place by default.
  - Some have argument `inplace`, and you may do  
`df.method_name(..., inplace=True)` instead of `df = df.method_name(...)`
- keywords:
  - axis: `0` or `index` -> along index; `1` or `columns` -> along columns
  - types: `str` , `category` , `float` , `int`

## Exercise 2

1. Download `british-english_a-e` and `american-english_a-e` file from absalon
2. Read them into two `DataFrame` variables `df1` and `df2`. Set the column name to `words` for both dataframes
3. Merge the two word lists to a dataframe `df`. Hint: check the argument `how` in `pd.merge()` (look in the documentation or search the web)

# What is Data Science?

**Data Science** is the intersection of:

- Statistics & Programming
- Domain Knowledge

**Key activities:**

- Data Collection and CLEANING
- Analysis
- Visualization
- Interpretation

# Data Science Process

## Data Collection

- Gather data that can help the analysis and manipulation

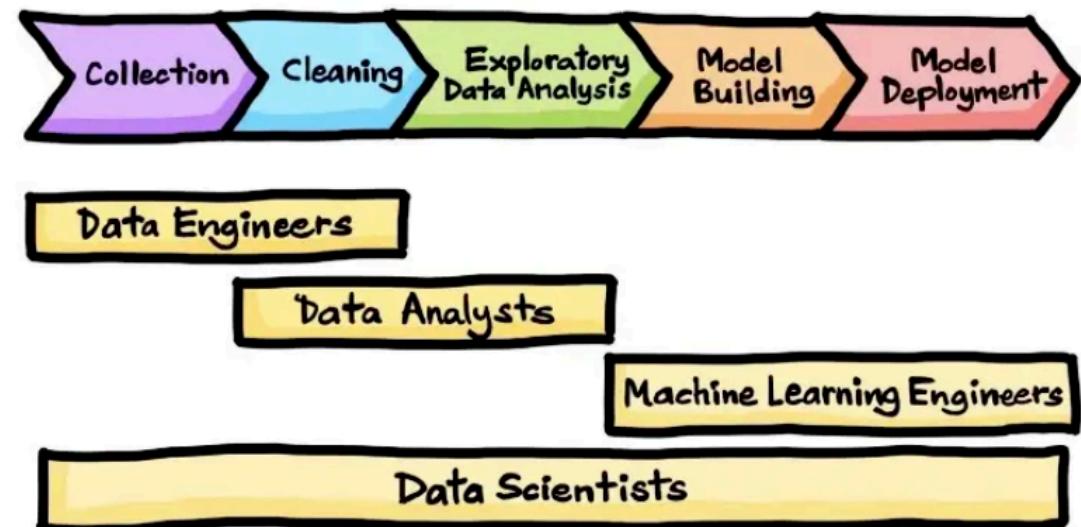
## Cleaning

- Clean the raw data and convert it into the desired form format

## Exploratory Data Analysis (EDA)

- Analyze the data, identify the relevant factors, suggests directions towards modeling

# THE DATA SCIENCE PROCESS



# Why NumPy in Data Science?

- **Efficient Array Operations**
  - Fast computation on large datasets
  - Memory-efficient storage
  - Vectorized operations (no explicit loops needed)
- **Mathematical Operations**
  - Linear algebra operations
  - Statistical functions
  - Random number generation

## NumPy Generating Random Data

```
np.random.rand(3,3)
array([[0.45537793, 0.8115471 , 0.74602629],
       [0.84160578, 0.36393521, 0.51246094],
       [0.34451085, 0.39302559, 0.88359236]])
```

```
np.random.standard_normal((3,3))
array([-0.85249074, -0.40046098, -0.58350246],
      [ 3.24900664, -0.0125803 , -0.34059938],
      [ 1.4437066 ,  0.34075023, -2.40908765]])
```

# NumPy Data Preprocessing

```
# Handle missing values
data = np.array([1, np.nan, 3, 4])
clean_data = np.nan_to_num(data) # replaces NaN with 0 or another defined value

# Normalize data
normalized = (data - np.mean(data)) / np.std(data)
```

# NumPy Feature Engineering

Useful for machine learning

```
# Create polynomial features
x = np.array([1, 2, 3, 4])
poly_features = np.column_stack([x, x**2, x**3])

array([[ 1,  1,  1],
       [ 2,  4,  8],
       [ 3,  9, 27],
       [ 4, 16, 64]])
```

# NumPy Statistical Analysis

```
# Basic statistics  
mean = np.mean(data)  
std = np.std(data)  
...
```

# Dataset Representation

- Rows: Samples/observations
- Columns: Features/variables

```
# Dataset with 3 samples, 4 features
dataset = np.array([
    [1.2, 2.3, 0.5, 1.0], # Sample 1
    [2.1, 1.5, 0.8, 2.2], # Sample 2
    [1.8, 1.9, 0.6, 1.5]  # Sample 3
])
```

# NumPy image Processing

- Images as 2D/3D arrays

```
# Grayscale image: 2D array
image = np.array([
    [255, 128, 0],
    [128, 255, 128],
    [0, 128, 255]
])
```

# Wrapping NumPy arrays in Pandas DataFrames or Series to get labeled data

You have seen it many time:

```
a = np.array([1, 2, 3])
df = pd.DataFrame({"values": a})
values
0      1
1      2
2      3
```

# How NumPy and Pandas Work Together

- Pandas uses NumPy arrays internally for data storage
- Different roles
  - NumPy: fast numerical engine
  - Pandas: structure, labels, and data handling
- Typical workflow
  - i. Use Pandas to load, clean, and organize data
  - ii. Use NumPy for fast numerical computations
  - iii. Store results back in a DataFrame

# NumPy vs Pandas

## Comparison Table

Feature	NumPy	Pandas
Main purpose	Numerical computing	Data analysis & manipulation
Core structures	ndarray	Series , DataFrame
Data types	Single data type per array	Mixed data types supported
Indexing	Integer-based	Label-based ( loc ) and integer ( iloc )
Typical use cases	Math, linear algebra, speed	Tables, CSVs, grouping, cleaning

# Data science exercises 2 and 3

Soccer and protein structure

Worksheet on Github and Absalon

Old exam exercises

**Still not confident with indexing and slicing arrays and DataFrames? Consider doing the Monday worksheet if you haven't already.**